

Machine Learning Worksheet 3

Thomas Blocher – MatrNr. 03624034

Raphael Dümig – MatrNr. 03623199

November 1, 2015

Problem 1

see addendum of jupyter notebook *kNN_implementationHW*

Problem 2

see addendum of jupyter notebook *decision-tree*

Problem 3

results from the decision tree implementation of problem 2:

$$\hat{z}_a = 1$$

$$\hat{z}_b = 2$$

$$p(c = \hat{z}_a | x_a, T) = 1$$

$$p(c = \hat{z}_b | x_b, T) = \frac{2}{3}$$

Problem 4

results from the 3NN implementation of problem 2 without Standardization:

$$\hat{z}_a = 0$$

$$\hat{z}_b = 2$$

results from the 3NN implementation of problem 2 with Standardization:

$$\hat{z}_a = 1$$

$$\hat{z}_b = 0$$

Problem 5

results from the 3NN implementation of problem 2 without Standardization regression:

$$\hat{z}_a = 1.00000$$

$$\hat{z}_b = 0.50904$$

Problem 6

LOOCV on the Dataset showed that the quality of 3-NN is not very good. Using standardization improves the quality a little bit but it's still below 30%. (see Octave Source code.) The decision trees perform better on the dataset.

Addendum:

kNN_implementationHW

October 31, 2015

1 Implementation exercise: KNN

Thomas Blocher (MatrNr. 03624034) Raphael Dümig (MatrNr. 03623199)

```
In [2]: import random
import numpy as np
import operator
from sklearn import datasets
import matplotlib.pyplot as plt
%matplotlib inline
```

```
/usr/lib/python3.5/site-packages/sklearn/utils/fixes.py:64: DeprecationWarning: inspect.getargspec() is
if 'order' in inspect.getargspec(np.copy)[0]:
```

1.1 Load dataset

The iris data set (https://en.wikipedia.org/wiki/Iris_flower_data_set) it loaded by the function loadDataset.
Arguments:

- *split*: int Split rate between test and training set e.g. 0.67 corresponds to 1/3 test and 2/3 validation

Returns:

- *X*: list(array of length 4); Trainig data
- *Z*: list(int); Training labels
- *XT*: list(array of length 4); Test data
- *ZT*: list(int); Test labels

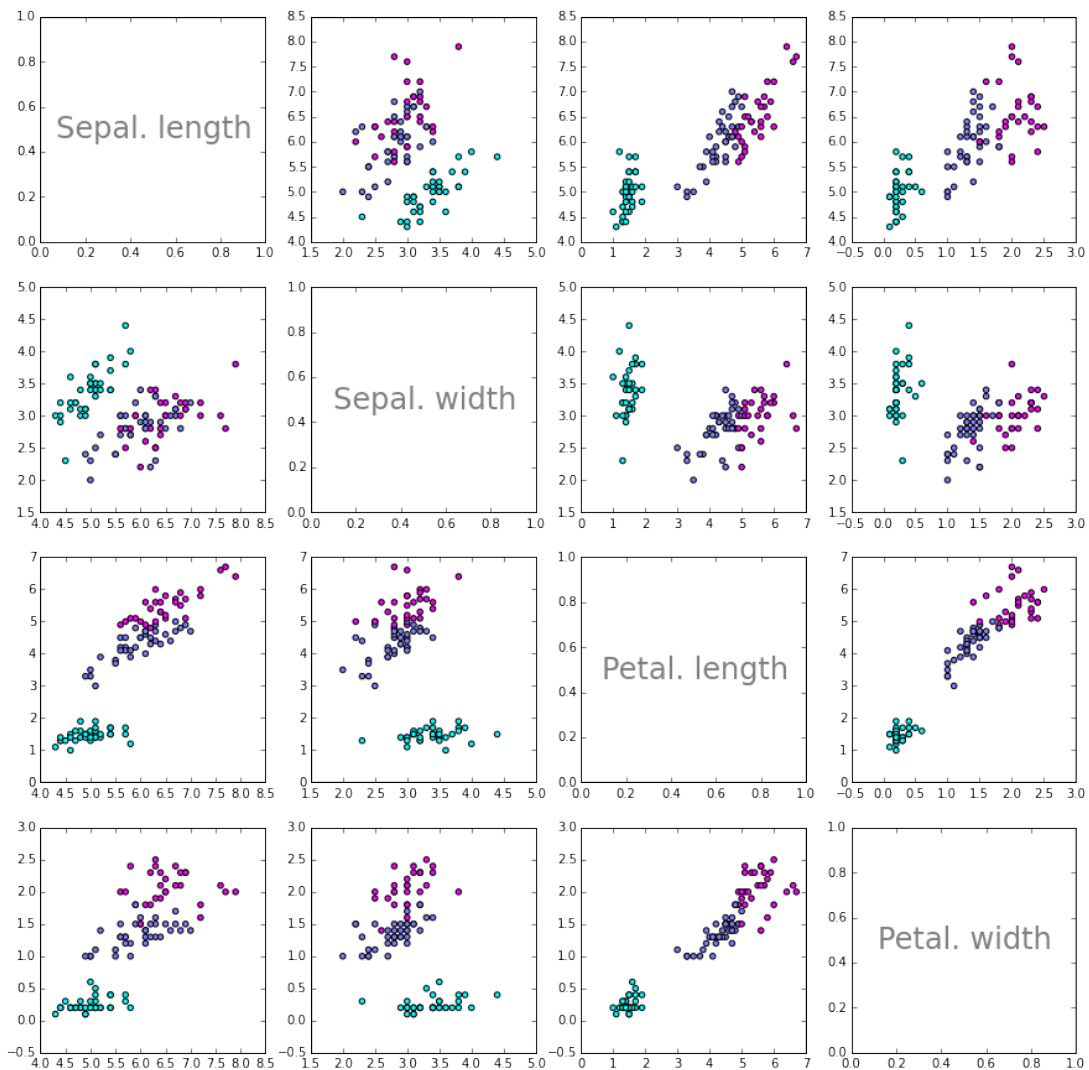
```
In [3]: def loadDataset(split, X=[], XT=[], Z = [], ZT = []):
dataset = datasets.load_iris()
c = list(zip(dataset['data'], dataset['target']))
random.seed(224)
random.shuffle(c)
x, t = zip(*c)
sp = int(split*len(c))
X = x[:sp]
XT = x[sp:]
Z = t[:sp]
ZT = t[sp:]
return X, XT, Z, ZT
```

```
In [4]: # prepare data
split = 0.67
X, XT, Z, ZT = loadDataset(split)
```

1.2 Plot dataset

Since X is dimensionality 4, 16 scatterplots (4x4) are plotted showing the dependencies of each two features.

```
In [5]: Xa = np.asarray(X)
f, axes = plt.subplots(4, 4, figsize=(15, 15))
for i in range(4):
    for j in range(4):
        if j == 0 and i == 0:
            axes[i,j].text(0.5, 0.5, 'Sepal. length', ha='center', va='center', size=24, alpha=0.5)
        elif j == 1 and i == 1:
            axes[i,j].text(0.5, 0.5, 'Sepal. width', ha='center', va='center', size=24, alpha=0.5)
        elif j == 2 and i == 2:
            axes[i,j].text(0.5, 0.5, 'Petal. length', ha='center', va='center', size=24, alpha=0.5)
        elif j == 3 and i == 3:
            axes[i,j].text(0.5, 0.5, 'Petal. width', ha='center', va='center', size=24, alpha=0.5)
        else:
            axes[i,j].scatter(Xa[:,j], Xa[:,i], c = Z, cmap=plt.cm.cool)
```



1.3 Exercise 1: Euclidean distance

Compute euclidean distance between two data points.

arguments: * *x1*: array of length 4; data point * *x2*: array of length 4; data point
returns: * *distance*:int; euclidean distance between *x1* and *x2*

```
In [7]: def euclideanDistance(x1, x2):  
        d = x1-x2  
        return np.sqrt(np.sum(d*d, axis=-1))
```

1.4 Exercise 2: get k nearest neighbours

For one data point *xt* compute all *k* nearest neighbours.

arguments: * *X*: list(array of length 4); Training data * *Z*: list(int); Training labels * *xt*: array of length 4; Test data point
returns: * neighbours: list of length *k* of tuples (*X_neighbor*, *Z_neighbor*, distance between neighbor and *xt*); **this is the list of k nearest neighbours to *xt***

```
In [10]: def getNeighbors(X, Z, xt, k):  
        dists = euclideanDistance(X, xt)  
  
        return [(X[i], Z[i], dists[i]) for i in np.argsort(dists)[:k]]
```

1.5 Exercise 3: get neighbor response

For the previously computed *k* nearest neighbors compute the actual response. I.e. give back the class of the majority of nearest neighbors. What do you do with a tie?

arguments: * neighbours
returns * *y*: int; majority target

```
In [16]: from collections import defaultdict  
  
def getResponse(neighbors):  
    l = {}  
    for n in neighbors:  
        try:  
            l[n[1]] += 1  
        except KeyError:  
            l[n[1]] = 1  
  
    n_max = max(l.values())  
    max_labels = list(filter(lambda k: l[k] == n_max, l.keys()))  
  
    return max_labels[0]
```

1.6 Exercise 4: Compute accuracy

arguments: * *YT*:list(int); predicted targets * *ZT*:list(int); actual targets
returns: * accuracy (percentage of correctly classified test data points)

```
In [12]: def getAccuracy(YT, ZT):  
        return np.sum(np.asarray(YT) == np.asarray(ZT)) / len(YT)
```

```
In [13]: def predict(X, Z, XT, k):
          Y=[]
          for xt in XT:
              neighbors = getNeighbors(X, Z, xt, k)
              Y.append(getResponse(neighbors))
          return Y
```

1.7 Testing

Should output an accuracy of 0.9599999999999996%.

```
In [17]: # prepare data
          split = 0.67
          X, XT, Z, ZT = loadDataset(split)
          print('Train set: ' + repr(len(X)))
          print('Test set: ' + repr(len(XT)))
          # generate predictions
          k = 3
          YT = predict(X, Z, XT, k)
          accuracy = getAccuracy(YT, ZT)
          print('Accuracy: ' + repr(accuracy) + '%')
```

```
Train set: 100
Test set: 50
Accuracy: 0.9599999999999996%
```

```
In [ ]:
```

decision_tree

October 31, 2015

```
In [3]: # Raphael Dümig (MatrNr. 03623199)
        # Thomas Blocher (MatrNr. 03624034)

        # load the data

import numpy as np
import matplotlib.pyplot as plt
import matplotlib.cm as cm
import math as m
%matplotlib inline

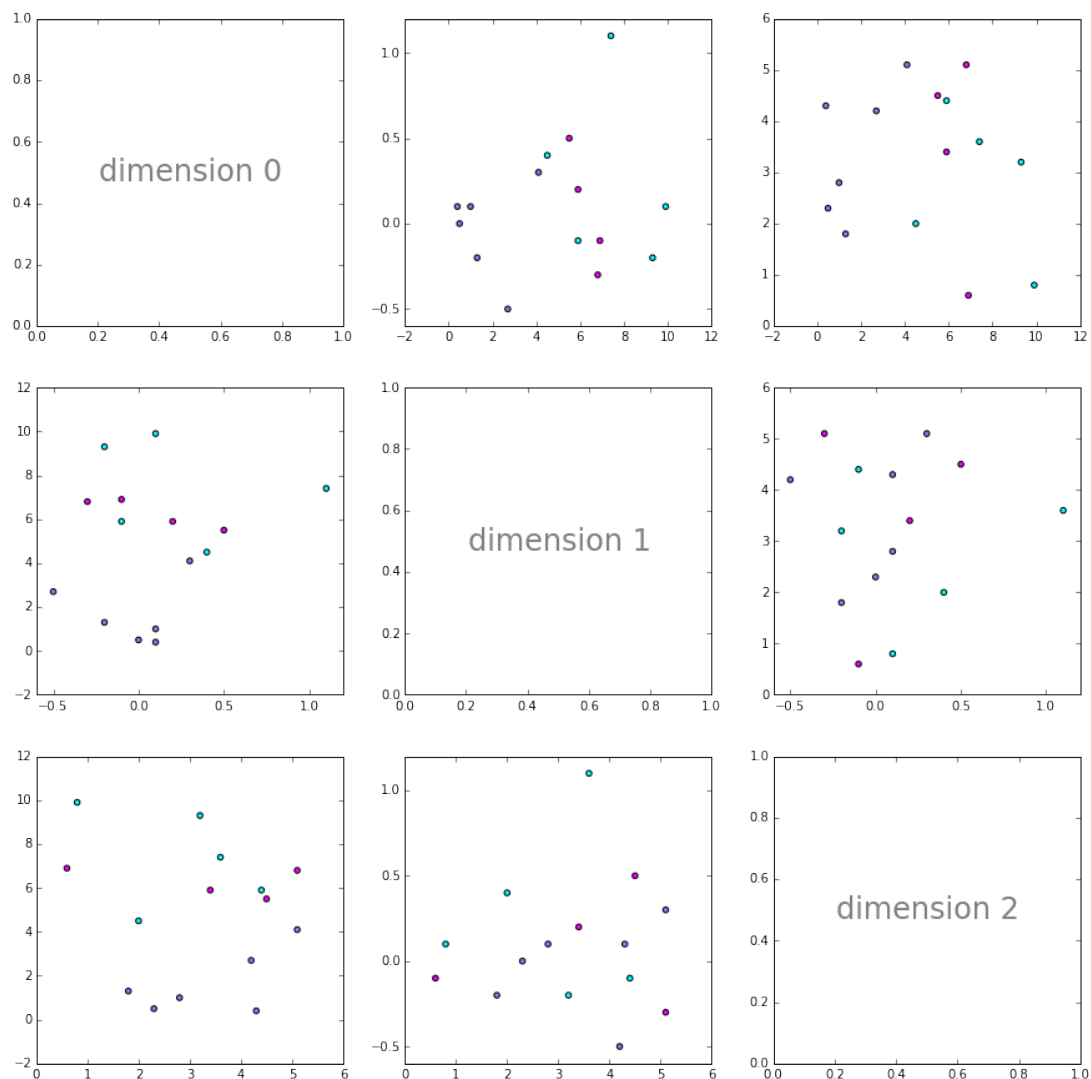
data = np.genfromtxt('../homework03.csv', delimiter=',', skip_header=1)

X = data[:, :3]
Z = data[:, 3]

In [10]: # plot the data

fig, axes = plt.subplots(3,3, figsize=(15,15))

for i in range(3):
    for j in range(3):
        if i != j:
            axes[i,j].scatter(X[:,i], X[:,j], c=Z, cmap=cm.cool)
        else:
            axes[i,j].text(0.5, 0.5, 'dimension %d' % i,
                           ha='center', va='center', size=24, alpha=.5)
```



In [4]: # build a decision tree using the data

```
import itertools
```

```
def gini_index(labels, label_values):
```

```
    if labels.size == 0:
```

```
        return 0
```

```
    else:
```

```
        n_label = np.sum(labels[:, np.newaxis] == label_values, axis=0)
```

```
        return 1 - np.sum(n_label**2)/np.sum(n_label)**2
```

```
class DecisionTreeNode:
```

```
    def __init__(self, X, Z, depth=np.infty):
```

```
        #print('creating node with a dataset of size %d (%d levels remaining)' % (X.shape[0], d
```

```
        self.X = X
```



```

        self.Z = Z
        self.left = None
        self.right = None
        self.split_pos = None
        self.n = Z.shape[0]
        self.depth = depth
        self.split()

def gini(self):
    if self.right is None:
        return gini_index(self.Z, list(set(self.Z)))
    else:
        return (self.right.gini()*self.right.n + self.left.gini()*self.left.n) / self.n

def get_cut_positions(self):
    s = np.sort( self.X, axis=0 )
    return (s[1:] + s[:-1]) / 2

def get_labels(self):
    return list(set(self.Z))

def is_pure(self):
    return len(self.get_labels()) == 1

def get_label_counts(self):
    ls = self.get_labels()
    n_label = np.sum(self.Z[:, np.newaxis] == ls, axis=0)
    return { l:n for (l,n) in zip(ls,n_label) }

def split(self):
    if self.is_pure():
        print('node is pure! refusing to split...')
        return
    if self.depth <= 0:
        print('reached maximum depth! refusing to split...')
        return

    label_values = self.get_labels()
    cuts = self.get_cut_positions()
    zz = cuts[:, :, np.newaxis] < self.X.T
    indices = np.array([
        [ ( gini_index(self.Z[ zz[i,j,:]], label_values)*np.sum( zz[i,j,:])
          + gini_index(self.Z[~zz[i,j,:]], label_values)*np.sum(~zz[i,j,:])
          )
          for j in range(zz.shape[1])
        ]
        for i in range(zz.shape[0])
    ])
    min_index = np.argmin(indices)
    split_dim = min_index % indices.shape[1]
    split_dest = min_index // indices.shape[1]
    self.split_pos = (split_dim , cuts[split_dest, split_dim])
    mask = zz[split_dest, split_dim, :]
    # put all elements greater than the splitting boundary into the left node

```

```

        self.left = DecisionTreeNode(self.X[mask], self.Z[mask], self.depth-1)
        # all others in the right node
        self.right = DecisionTreeNode(self.X[~mask], self.Z[~mask], self.depth-1)
        return

    def subtree_to_str(self, d=0):
        res = self.node_to_str(d)
        if self.left is not None:
            res += '\n'.join(['|' + l for l in itertools.chain(
                self.left.subtree_to_str(d+1).split('\n'),
                self.right.subtree_to_str(d+1).split('\n'))
            ])

        return res

    def node_to_str(self, d=0):
        res = ''
        if self.split_pos is not None:
            res = ('depth: %d -- X%s > %.2f' % (d, chr(ord('0') + self.split_pos[0] + 1),
                                                self.split_pos[1]))
        else:
            res = ('depth: %d -- leaf' % d)
        res += ' (gini: %.2f)\n' % self.gini()
        nl = self.get_label_counts()
        for l in sorted(nl.keys()):
            res += ('\t%s: %f%%\n' % (str(l), 100*nl[l]/self.n))
        return res

class DecisionTree:
    def __init__(self, X, Z, max_depth=np.infty):
        self.X = X
        self.Z = Z
        self.max_depth = max_depth
        self.root = DecisionTreeNode(X,Z,max_depth)

    def gini_index(self):
        return self.root.gini_index()

    def print_tree(self):
        print('DFS traversal of the tree:')
        print(self.root.subtree_to_str(0))

In [5]: T = DecisionTree(X,Z.astype(int),2)
        print(30 * '=' + '\n')
        T.print_tree()

node is pure! refusing to split...
reached maximum depth! refusing to split...
node is pure! refusing to split...
=====

DFS traversal of the tree:
depth: 0 -- X1 > 4.30 (gini: 0.18)
      0: 33.333333%

```

```

      1: 40.000000%
      2: 26.666667%
|depth: 1 --  $X_1 > 7.15$  (gini: 0.30)
|      0: 55.555556%
|      2: 44.444444%
||depth: 2 -- leaf (gini: 0.00)
||      0: 100.000000%
||
||depth: 2 -- leaf (gini: 0.44)
||      0: 33.333333%
||      2: 66.666667%
||
|depth: 1 -- leaf (gini: 0.00)
|      1: 100.000000%
|

```

Source code: 3NN

```
1 %% Read in Dataset
2 data = csvread('homework03.csv');
3 %remove first line (Remove labels)
4 data = data(2:end,:);
5 coords = data(:,1:3)
6 labels = data(:,4)
7
8 xa = [4.1, -0.1, 2.2]
9 xb = [6.1, 0.4, 1.3]
10
11
12 %% Since the axis are not equally derivated we have to do some
    standardization
13
14 %% Doing the Standardization
15 xaStd(1) = (xa(1) - mean(coords(:,1)))./std(coords(:,1));
16 xaStd(2) = (xa(2) - mean(coords(:,2)))./std(coords(:,2));
17 xaStd(3) = (xa(3) - mean(coords(:,3)))./std(coords(:,3));
18
19 xbStd(1) = (xb(1) - mean(coords(:,1)))./std(coords(:,1));
20 xbStd(2) = (xb(2) - mean(coords(:,2)))./std(coords(:,2));
21 xbStd(3) = (xb(3) - mean(coords(:,3)))./std(coords(:,3));
22
23 coordsStd(:,1) = (coords(:,1) - mean(coords(:,1)))./std(coords
    (:,1));
24 coordsStd(:,2) = (coords(:,2) - mean(coords(:,2)))./std(coords
    (:,2));
25 coordsStd(:,3) = (coords(:,3) - mean(coords(:,3)))./std(coords
    (:,3));
26 %%LOOCV Tests
27 %Without Standardization
28 for valid = 1:size(labels,1)
29     dist = [];
30     x = data(valid,1:3);
31     for i = 1:size(labels,1)
32         %Calculating the dist to each point
33         dist = [norm(coords(i)-x),dist];
34     end
35     [~,i] = sort(dist);
36     %Since Training data are unique we get the 3 NN to Validation
        Point by taking the 3 Closest point with dist ~= 0 (i(1)=
        valid)
37     % the second part of the expression returns true iff there is a
        tie
38     out(valid) = (mode(labels(i(2:4))) == labels(valid)) || size(
        unique(labels(i(2:4))),2) == 3;
```

```

39 end
40
41 %With Standardization
42 for valid = 1:size(labels,1)
43     dist = [];
44     x = coords(valid,1:3);
45     for i = 1:size(labels,1)
46         dist = [norm(coordsStd(i,:)-x), dist];
47     end
48     [~,i] = sort(dist);
49     outStd(valid) = (mode(labels(i(2:4))) == labels(valid)) || size
        (unique(labels(i(2:4))),2) == 3;
50 end
51
52 %%%LOOCV Results
53
54 quality = mean(out)
55
56 qualityStd = mean(outStd)
57
58 %% Labeling the Points
59 %Classification without Standardization
60 for i = 1:size(labels,1)
61     dist(i) = norm(coords(i,:)-xa);
62 end
63 [~,i] = sort(dist);
64 va = mode(data(i(1:3),4))
65
66 for i = 1:size(labels,1)
67     dist(i) = norm(coords(i,:)-xb);
68 end
69
70 [~,index] = sort(dist);
71 vb = mode(data(index(1:3),4))
72
73
74
75 %Classification with Standardization
76 for i = 1:size(labels,1)
77     dist(i) = norm(coordsStd(i,:)-xaStd);
78 end
79
80 [~,i] = sort(dist);
81 vaStandardized = mode(labels(i(1:3)))
82
83 for i = 1:size(labels,1)
84     dist(i) = norm(coordsStd(i,:)-xbStd);
85 end
86

```

```

87 [~,i] = sort(dist);
88 vbStandardized = mode(labels(i(1:3)))
89
90
91 %% Regression (just using standardization)
92
93 for i = 1:size(labels,1)
94     dist(i) = norm(coordsStd(i,:) - xaStd);
95 end
96
97 [~,i] = sort(dist);
98 vareg = (1/sum(1./dist(i(1:3)))) * sum ( 1./(dist(i(1:3))) *
99     labels(i(1:3)))
100
101 for i = 1:size(labels,1)
102     dist(i) = norm(coordsStd(i,1:3) - xbStd);
103 end
104
105 [~,i] = sort(dist);
106 vbreg = (1/sum(1./dist(i(1:3)))) * sum ( 1./(dist(i(1:3))) *
107     labels(i(1:3)))

```

Execution results

```

octave>hw2_4_5
coords =

```

5.50000	0.50000	4.50000
7.40000	1.10000	3.60000
5.90000	0.20000	3.40000
9.90000	0.10000	0.80000
6.90000	-0.10000	0.60000
6.80000	-0.30000	5.10000
4.10000	0.30000	5.10000
1.30000	-0.20000	1.80000
4.50000	0.40000	2.00000
0.50000	0.00000	2.30000
5.90000	-0.10000	4.40000
9.30000	-0.20000	3.20000
1.00000	0.10000	2.80000
0.40000	0.10000	4.30000
2.70000	-0.50000	4.20000

```

labels =

```

```

2
0

```

```

2
0
2
2
1
1
0
1
0
0
1
1
1

xa =

    4.10000   -0.10000    2.20000

xb =

    6.10000    0.40000    1.30000

quality = 0.13333
qualityStd = 0.26667
va = 0
vb = 2
vaStandardized = 1
vbStandardized = 0
vareg = 1.00000
vbreg = 0.50904

```