

Machine Learning Worksheet 04

Maximum Likelihood Estimation

1 Still refreshing ...

Problem 1: An unbiased coin is flipped until one head is thrown. What is the expected number of tails T and the expected number of heads H ? Show your work.

Problem 2: There are eleven urns labeled by $u \in \{0, 1, 2, \dots, 10\}$, each containing ten balls. Urn u contains u black balls and $10 - u$ white balls. Alice selects an urn u at random and draws N times with replacement from that urn, obtaining n_B black balls and $N - n_B$ white balls. If after $N = 10$ draws $n_B = 3$ black balls have been drawn, what is the probability that the urn Alice is using is urn u ?

Now, let Alice draw another ball from the same urn. What is the probability that the next drawn ball is black (show your work)?

2 Parameter Estimation

Consider n samples x_1, \dots, x_n drawn independently and identically (i.i.d.) from a given distribution $P(X \mid \theta)$. This distribution is usually parametrised (e.g., one parameter representing its mean, one its variance, etc.); these parameters are denoted by θ . One wants to find accurate estimates for these parameters using the n samples only. *Maximum Likelihood Estimation* (MLE) finds estimates for the various parameters at hand by maximising the likelihood $P(x_1, x_2, \dots, x_n \mid \theta) = \prod_{i=1}^n P(x_i \mid \theta)$ (i.e., the probability of observing the n samples at hand). Note that usually one considers the *log likelihood*, $\log P(x_1, \dots, x_n \mid \theta)$.

2.1 Coins

Let X be a Bernoulli random variable. The Bernoulli distribution is only parametrised by one parameter, $\theta = P(X = 1)$.

Problem 3: For n i.i.d. observations of X determine the MLE for θ . You might want to use $P(X = x \mid \theta) = \theta^x (1 - \theta)^{1-x}$.

Now we look at a slightly more complex distribution, the binomial distribution.

Problem 4: Consider a Bernoulli random variable X and suppose we have observed m occurrences of $X = 1$ and l occurrences of $X = 0$ in a sequence of $m + l$ Bernoulli experiments. We are only interested in the number of occurrences of $X = 1$ —we will model this with a Binomial distribution with parameter μ . A prior distribution for μ is given by the Beta distribution. Show that the posterior *mean* value of

μ lies between the prior mean of μ and the maximum likelihood estimate for μ . To do this, show that the posterior mean can be written as λ times the prior mean plus $(1 - \lambda)$ times the maximum likelihood estimate, with $0 \leq \lambda \leq 1$. This illustrates the concept of the posterior mean being a compromise between the prior distribution and the maximum likelihood solution.

Note: The probability mass function for the Binomial distribution is defined as follows:

$$p(x = m \mid N, \mu) = \binom{N}{m} \mu^m (1 - \mu)^{N-m}.$$

2.2 Poisson distribution

Let X be Poisson distributed.

Problem 5: Again, for n i.i.d. samples from X , determine the maximum likelihood estimate for λ . Show that this estimate is unbiased!

In class we also talked about avoiding overfitting of parameters via *prior* information. Compute the posterior distribution over λ , assuming a $Gamma(\alpha, \beta)$ prior for it. Compute the MAP for λ under this prior. Show your work.